

Skin Cancer Detection

CSCE 485 Computer and Machine Vision

Christopher Yang

November 18, 2024

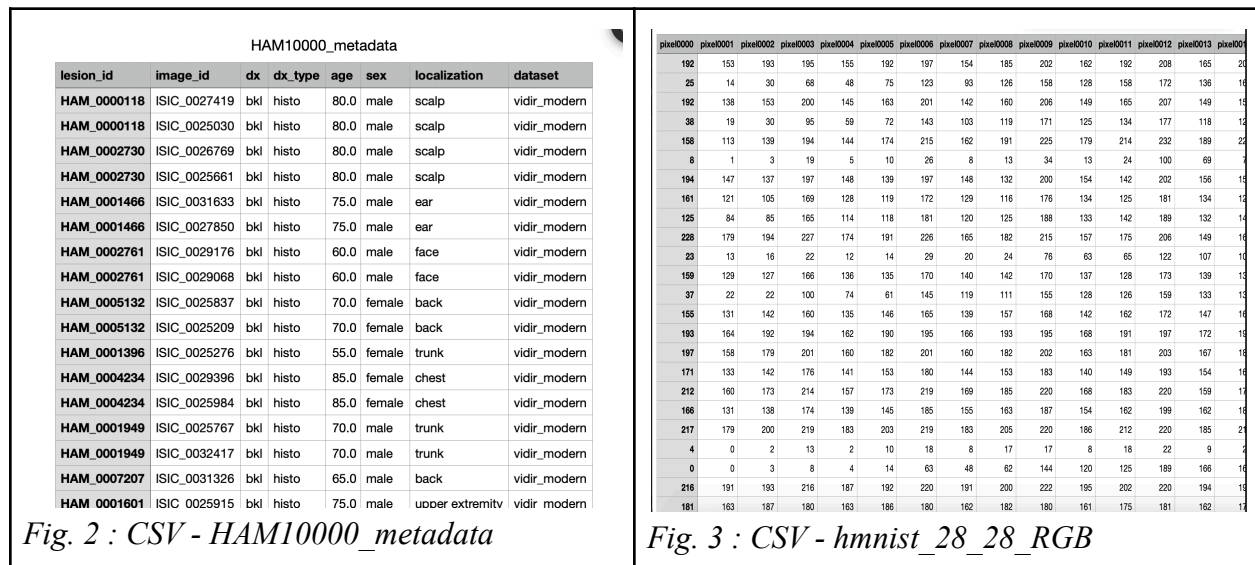
Introduction:

So the topic I chose for this project was skin cancer detection. The reason why I chose this topic is because skin cancer often goes undetected by many people, and healthcare professionals sometimes miss early signs, leading to delayed treatment and, in some cases, preventable death. Thus, I wanted to try to help make skin cancer more detectable in early stages so as to prevent preventable deaths.

Dataset:

The dataset I chose to work with was the HAM10000 dataset. This dataset has approximately 2.9 GB of 10,015 images of various types of skin cancer. The dataset also contained information on the location where the skin cancer was found, and the gender of each individual.






They also provide .csv files on the labeling and the pixel sizes of each sample. For figure 2, that CSV contains information on their patients, such as: age, gender, type of skin cancer, and location of the skin cancer. In figure 3, each image's data is stored as a series of numbers, where each number represents the color value (Red, Green, and Blue) for each pixel in the image.

Data Augmentation:

For the data, I applied several preprocessing steps for model training and to improve its generalization ability. First, I resized all the images to a uniform size of 28x28 pixels, ensuring consistency in the input data. Next, I normalized the images by scaling the pixel values between 0 and 1, which helps the model learn more efficiently by reducing the impact of large differences in feature scales. After normalization, I flattened the images from their original 2D structure into 1D vectors, which is necessary for the fully connected layers of the neural network. The reason why I did this is to help the model generalize better and avoid overfitting.



Methodology:

This project uses the HAM10000 dataset, consisting of 10,015 images of skin lesions classified into 7 different types of skin cancer. The images are prepared for the model by using the `hmnist_28_28_RGB.csv` file, which contains the pixel values of images that have been resized to 28x28 pixels. These pixel values are used in the model. The images are then normalized by scaling the pixel values to the range of 0 to 1, which helps improve training efficiency and stability.

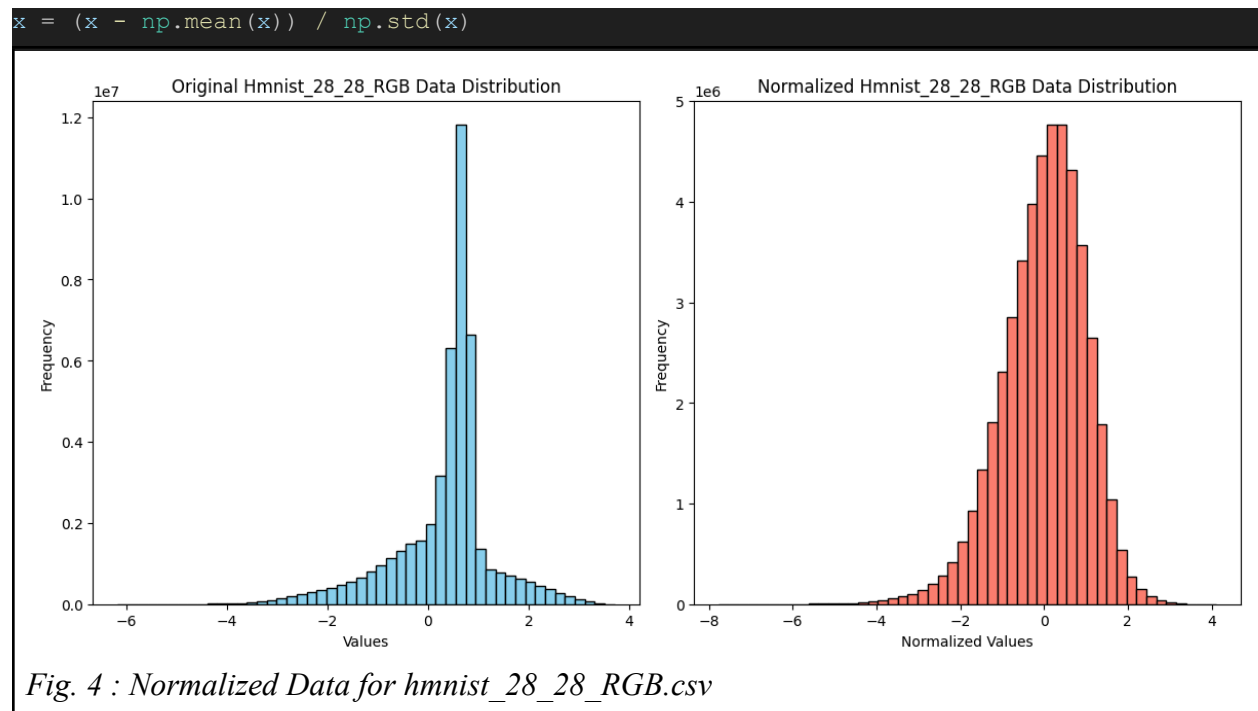
The model architecture has a fully connected neural network (MLP) designed to classify the images. The model starts with an `InputLayer`, where the input shape matches the 28x28 pixel images. After flattening the image data into a 1D vector, the model passes the data through a Dense layer with 128 units and ReLU activation. Finally, the model ends with a Dense layer of 7 units, to correspond for the seven class, and a softmax activation function.

I then used the categorical crossentropy loss function, used for multi-class classification, and the Adam optimizer, for efficiency in training neural networks. The model's performance is evaluated using a confusion matrix, which helps to visualize how well the model is distinguishing between different types of skin cancer. And I will visualize eight randomly



selected images alongside their true labels and predicted labels in a grid. To visually evaluate how well the model is performing.

Initial Results:

Before running the actual model, I normalize the dataset from figure 4 (right graph) so as to get an even distribution of my data and keep it in the same range. As in figure 4 (left graph), I'm dealing with a negatively skewed graph, which can lead to inaccurate predictions. So to solve that issue, I normalized the dataset to be a bell shape, as it makes the data more balanced, which could help with better predictions.



Afterwards, once I test and train the model with the dataset. I got the model for the accuracy and loss of the model here:



During the training and validation test, from the model, there was an accuracy of 67.89% and a loss of 90.27%. In figures 6 and 7, it appears that overfitting is occurring because the training accuracy remains significantly higher than the validation accuracy, and the validation loss increases as the model continues training. This indicates that the model is performing well on the training data but is failing to generalize to unseen data.

Confusion matrix:

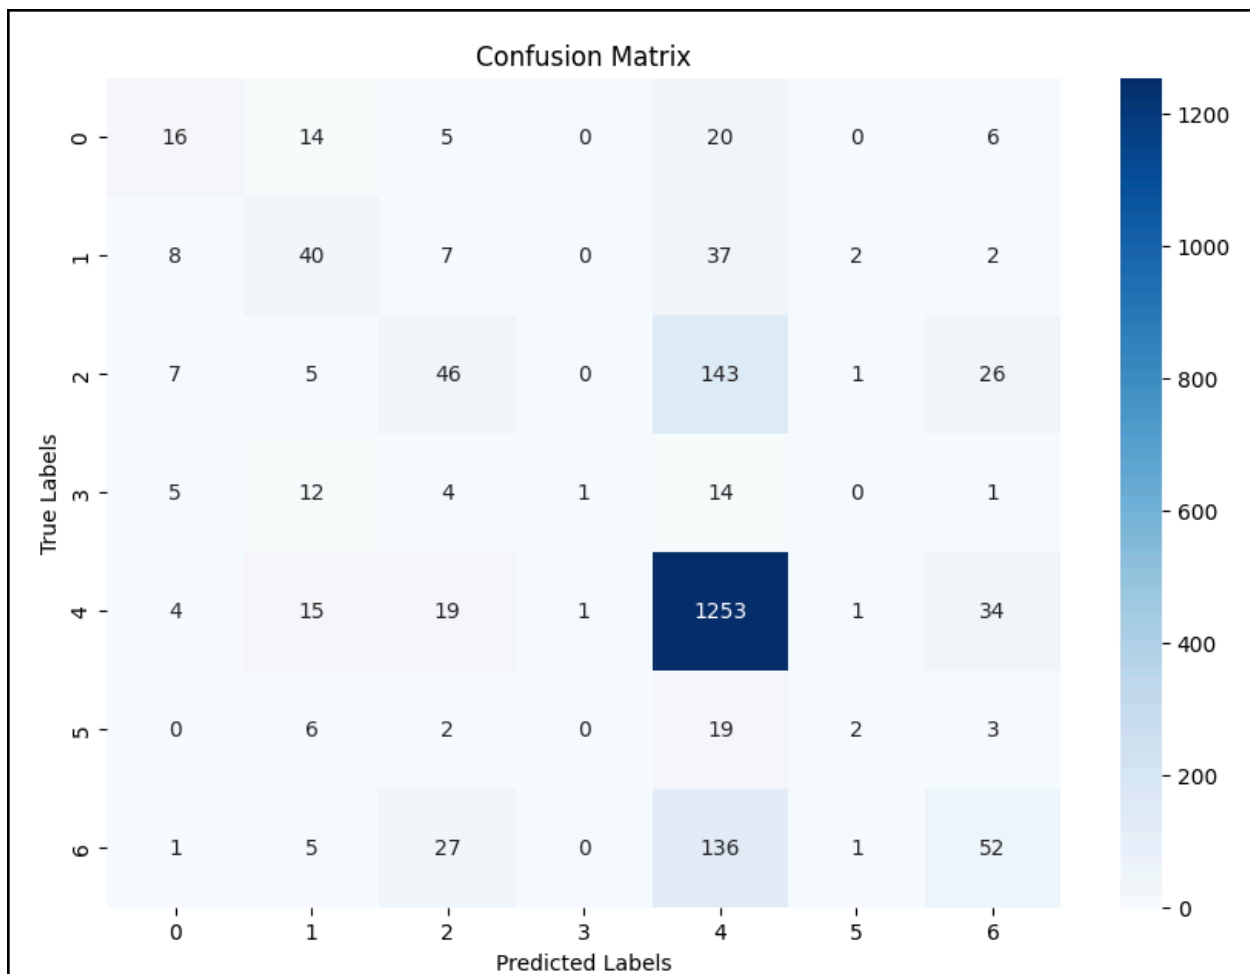


Fig. 8 : Confusion Matrix (Accuracy 67.89%, Loss 90.27%)

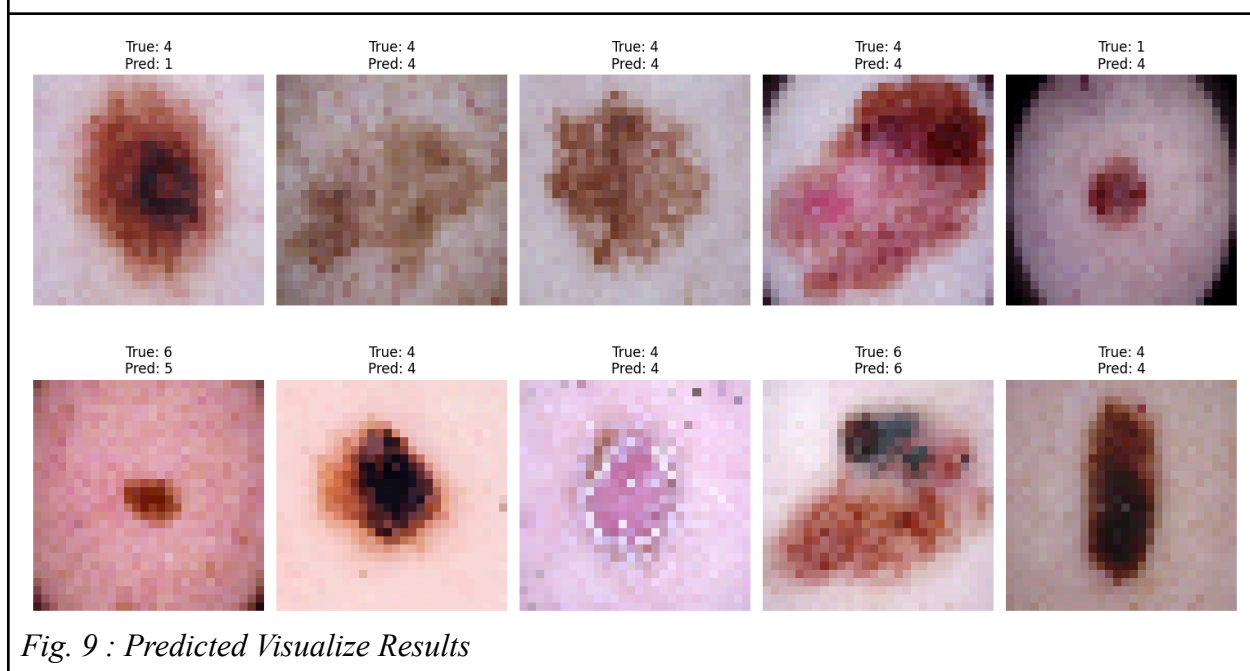


Fig. 9 : Predicted Visualize Results


Labels/Classes:

```
0: ('akiec', 'Actinic keratoses and intraepithelial carcinomae'),  
1: ('bcc', 'basal cell carcinoma'),  
2: ('bkl', 'benign keratosis-like lesions'),  
3: ('df', 'dermatofibroma'),  
4: ('nv', 'melanocytic nevi'),  
5: ('vasc', 'pyogenic granulomas and hemorrhage'),  
6: ('mel', 'melanoma')
```

Fig. 10 : Table of context

The first thing I noticed was that the prediction and true values in figure 9, were often predicted correctly. But for some of the cases when I reran the cell to get a different sample prediction, it would always mistake the sample as 4 (nv), which explains the reason why I got a loss of 90.27%. But when I made a confusion matrix, figure 8, I realized that my dataset is very imbalanced. As almost all of the skin cancer samples are being classified as melanocytic nevi (nv). So to make sure, I was dealing with an imbalancing problem. I used MATLAB and plotted the distribution of my dataset to see if I was indeed dealing with a balancing issue.

Attempts for Improvement (methods used are **bolded**):


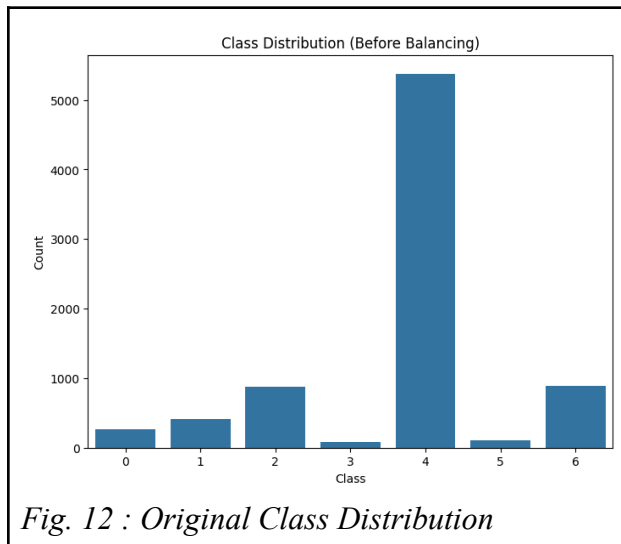


From looking at the confusion matrix and the distribution graph, it made it pretty obvious I was dealing with an imbalancing issue with my dataset. Because I was getting an accuracy of about 67.89%, I thought my model did well. But because my class distribution is so skewed. As there are more samples of nv than the rest, it explains the reasoning why I was getting 67.89%.

So from my understanding and inference, since the majority of the dataset belongs to the highest class distribution, such as the nv class, the model has a higher chance of accurately predicting the nv class. This is because nv samples make up approximately 75% of the dataset, which I believe explains why I was getting a high accuracy but an even higher loss.


To find a way to address the issue of class imbalance in my dataset, I did some research on techniques that could help. Initially, I was considering using the **class-weight** method like in class, which adjusts the weight of each class during training to make the model more sensitive to the minority class. However, after further consideration, I decided against using this method due to the fact that it didn't improve performance significantly in my tests. Whenever I added weights to the minority class, I would still have issues with overfitting and a very high loss in the 90%.

Instead, I decided to use **SMOTE** (Synthetic Minority Over-sampling Technique), a more advanced method of handling imbalanced datasets, it takes the minority class instances and generates synthetic data samples by interpolation, making the dataset more balanced without needing to collect new data.



To address that issue, I had with figure 4 and 5, I tried to reduce the **learning rate to 0.0001** and add a **dropout layer** in my model, with a **0.5 rate**, to help prevent overfitting. As a result there was a slight improvement in both models. As the accuracy became 69.90% and the loss was 90.16%, which isn't the best, but it was a slight improvement in the accuracy, but I was still dealing with a high loss of 90.16%. So to fix this issue with the high loss, I decided to add in the **early stop method** (results in "Final Results" section.)

For the models, I tried to implement another one, a basic **AlexNet model**, which had 5 convolution layers, 3 max-pooling layers, 2 Normalized layers, 2 fully connected layers and 1 SoftMax layer.



So, from the results above. I chose to not use that model for this project, because the accuracy was as low as 15%, which I think a factor to this model not performing well is because of the class imbalancing I was facing. I **re-augmented the data, to 227x227x3 images**, and increased the **epoch to about 30**, but I was still facing a lower accuracy of 12% and a high loss of about 110%.

Finally, the last method, I used was a library: `from imblearn.over_sampling import RandomOverSampler`

I tried using this library to solve the imbalancing issue; the library is used for oversampling the minority class in imbalanced datasets. So I thought this library would work really well for my model, but the only thing that it was able to work well on was only in solving the imbalancing issue I was having. The loss and accuracy models were very bad; there was a lot of overfitting, which I think is because of the library, because it randomly duplicates the samples in different classes, which I think increased the likelihood of overfitting since it just duplicates existing data rather than creating new synthetic examples. Because of that, I choose to stick with SMOTE which creates synthetic examples through interpolation. Note, this method did not improve the results for the AlexNet model also.

Final Results:




Fig. 14 : Final Model Accuracy




Fig. 15 : Final Model Loss

So for my final results, I stuck with my original model that I described in the methodology. I used SMOTE to solve the issue, where I had an imbalancing issue and for the overfitting issue, I used the early stop method. So when testing with the model, I tested different epochs, ranging from 10-50, but 10 performed much better. As I was getting an accuracy of 80 - 96% and a loss of 10-21 % which was much better than my original result. Notably, the confusion matrix was now predicting the other classes about 80 - 90% of the time from what I inferred in figure 16 and 17.






Fig. 16 : Final Confusion Matrix (Accuracy 94.43%, Loss 19.73%)





What Would I Do Differently:

What I would do differently, would be trying to get the AlexNet model to work. I spent too long trying to solve the imbalancing issue and overfitting. I just ran out of time to figure out what was going on with the AlexNet model. So, maybe next time, I would work with this model first, because it's known to work well with big datasets.

Conclusion:

Overall, I learned a lot about how to preprocess, analyze, and work with imbalanced datasets. I learned how crucial it is to properly prepare data for machine learning models by scaling images to a consistent dimension and normalizing datasets. Also, of the difficulties in dealing with deep learning architectures, and how important data balancing is for better model performance.